


P0 locked baseline: assignment evolution (120/120 epochs)



One training run; all augmented training batches. No extrapolation or uncertainty estimates.